

Autoencoders

Vahid Tarokh

ECE 685D, Fall 2025

Unsupervised Learning


- Unsupervised learning: we only use the inputs $\mathbf{x}^{(t)}$ for learning
 - automatically extract meaningful features for your data
 - leverage the availability of unlabeled data
 - add a data-dependent regularizer to training (in some cases)
- We will consider 3 models for unsupervised learning that will form the basic building blocks for deeper models:
 - **Autoencoders**
 - Sparse Coding
 - Restricted Boltzmann Machines

Autoencoders

- Map high-dimensional data to lower dimensions
 - Visualization
 - Compression (reducing the file size)
- Learn abstract features in an unsupervised way for downstream supervised tasks
 - Unlabeled data usually are much more plentiful than labeled data
- Build some generative models

Autoencoders

An autoencoder is a feed-forward neural net whose job is to take an input x and output \hat{x}




- Details of what goes inside the encoder and decoder matter!

Totally Linear Case: PCA

- The simplest kind of autoencoder has one hidden layer, linear activations, and squared error loss.

$$\mathcal{L}(\mathbf{x}, \tilde{\mathbf{x}}) = \|\mathbf{x} - \tilde{\mathbf{x}}\|^2$$

- This network computes $\tilde{\mathbf{x}} = \mathbf{U}\mathbf{V}\mathbf{x}$, which is a linear function.
- If $K \geq d$, we can choose \mathbf{U} and \mathbf{V} such that $\mathbf{U}\mathbf{V}$ is the identity. This isn't very interesting.



If $k < d$, then we have dimensionality reduction.

A linear autoencoder with MSE loss, one hidden layer, and tied weights learns the exact same subspace as PCA.
Autoencoders are a nonlinear generalization of PCA.

PCA

- Dimensionality reduction technique for data in \mathbb{R}^d
- Problem statement
 - ▶ given m points x_1, \dots, x_m in \mathbb{R}^d
 - ▶ and target dimension $k < d$
 - ▶ find “best” k -dimensional subspace approximating the data
- Formally: find matrices $U \in \mathbb{R}^{d \times k}$ and $V \in \mathbb{R}^{k \times d}$
- that minimize

$$f(U, V) = \sum_{i=1}^m \|x_i - UVx_i\|_2^2$$

- $V : \mathbb{R}^d \rightarrow \mathbb{R}^k$ is “compressor”, $U : \mathbb{R}^k \rightarrow \mathbb{R}^d$ is “decompressor”
- Is $f : \mathbb{R}^{d \times k} \times \mathbb{R}^{k \times d} \rightarrow \mathbb{R}$ convex?
- No! $f(U, \cdot)$ and $f(\cdot, V)$ both convex, but not $f(\cdot, \cdot)$
- **Claim:** optimal solution achieved at $U = V^\top$ and $U^\top U = I$
(columns of U are orthonormal)


PCA

- Optimization problem: minimize $_{U \in \mathbb{R}^{d \times k}, V \in \mathbb{R}^{k \times d}} \sum_{i=1}^m \|x_i - UVx_i\|_2^2$
- **Claim:** optimal solution achieved at $U = V^\top$ and $U^\top U = I$
- Proof:
 - ▶ For any U, V , linear map $x \mapsto UVx$ has range R of dimension k
 - ▶ Let w_1, \dots, w_k be orthonormal basis for R ; arrange into columns of W .
 - ▶ Hence, for each x_i there is $z_i \in \mathbb{R}^k$ such that $UVx_i = Wz_i$.
 - ▶ Note: $W^\top W = I$.
 - ▶ Which z minimizes $f(x_i, z) := \|x_i - Wz\|_2^2$?
 - ▶ For all $x \in \mathbb{R}^d, z \in \mathbb{R}^k$,
$$f(x, z) = \|x\|_2^2 + z^\top W^\top Wz - 2z^\top W^\top x = \|x\|_2^2 + \|z\|_2^2 - 2z^\top W^\top x.$$
 - ▶ Minimize w.r.t. z : $\nabla_z f = 2z - 2W^\top x = 0 \implies z = W^\top x$.
 - ▶ Therefore

$$\sum_{i=1}^m \|x_i - UVx_i\|_2^2 = \sum_{i=1}^m \|x_i - Wz_i\|_2^2 \geq \sum_{i=1}^m \|x_i - WW^\top x_i\|_2^2.$$

- ▶ U, V are optimal, so $\sum_{i=1}^m \|x_i - UVx_i\|_2^2 = \sum_{i=1}^m \|x_i - WW^\top x_i\|_2^2$.
- ▶ So instead of U, V can take W, W^\top . \square
- $WW^\top x$ is the **orthogonal projection** of x onto R .

Linear Autoencoder (PCA)



- If the **hidden and output layers** are linear, it will learn hidden units that are a linear function of the data and minimize the squared error.
- The K hidden units will span the same space as the first k principal components. The weight vectors may not be orthogonal.
- If $K < d$, and W with orthogonal *rows*, then we have PCA.

With nonlinear hidden units, we have a **nonlinear generalization of PCA**.

Optimality of the Linear Autoencoder

- Let us consider the following theorem:
 - let \mathbf{A} be the empirical covariance of data points
 - Let $\mathbf{A} = \mathbf{U} \Sigma \mathbf{U}^\top$
 - singular value decomposition
 - Σ is a diagonal matrix
 - \mathbf{U} are orthonormal matrices (columns/rows are orthonormal vectors)

Linear Autoencoder (PCA)


- The optimal pair of encoder and decoder is given by:

$$\mathbf{h}(\mathbf{x}) = (\mathbf{U}_{\cdot, \leq k})^\top \mathbf{x}$$

$$\underbrace{\mathbf{W}}$$

$$\hat{\mathbf{x}} = \mathbf{U}_{\cdot, \leq k} \mathbf{h}(\mathbf{x})$$

$$\underbrace{\mathbf{W}^*}$$



- for the sum of squared difference error
- for an auto-encoder with a linear decoder
- where optimality means “has the lowest training reconstruction error”

Optimality of the Linear Autoencoder


- So an optimal pair of encoder and decoder is

$$\mathbf{h}(\mathbf{x}) = (\mathbf{U}_{\cdot, \leq k})^\top \mathbf{x}$$

$$\underbrace{\mathbf{W}}$$

$$\widehat{\mathbf{x}} = \mathbf{U}_{\cdot, \leq k} \mathbf{h}(\mathbf{x})$$

$$\underbrace{\mathbf{W}^*}$$




- If inputs are normalized as follows:

$$\mathbf{x}^{(t)} \leftarrow \frac{1}{\sqrt{T}} \left(\mathbf{x}^{(t)} - \frac{1}{T} \sum_{t'=1}^T \mathbf{x}^{(t')} \right)$$

- encoder corresponds to **Principal Component Analysis (PCA)**
- singular values and (left) vectors = the eigenvalues/vectors of covariance matrix

Example: Applications of PCA

- Face Recognition
- An image is a point in a high-dimensional space
 - An $N \times M$ image is a point in R^{NM}
 - We can define vectors in this space



- We can find the best subspace using PCA
- Suppose this is a K-dimensional subspace
- This is like fitting a “hyper-plane” to the set of faces
 - spanned by vectors v_1, v_2, \dots, v_K
 - any face $x \sim a_1v_1 + a_2v_2 + \dots + a_Kv_K$
- The set of faces is a “subspace” of the set of images.

Application of PCA

- Let F_1, F_2, \dots, F_M be a set of training face images.
- Let F be their mean and $f_i = F_i - F$
- Use principal components to compute the eigenvectors and eigenvalues of the covariance matrix of the f_i 's
- Choose the vector u of most significant M eigenvectors to use as the basis.
- Each face (with F removed) is represented as a linear combination of eigenfaces $u = (u_1, u_2, u_3, u_4, u_5)$;
- $f_{27} = a_1 * u_1 + a_2 * u_2 + \dots + a_5 * u_5$



Find the face class k that minimizes

$$\|A - A_k\|$$

Application of PCA

training
images



mean
image




3 eigen-
images

linear
approxi-
mations



Nonlinear Autoencoder (Example)

- Feed-forward neural network trained to reproduce its input at the output layer
- Nonlinear generalization of PCA



Decoder

$$\begin{aligned}\hat{x} &= \hat{a}(x) \\ &= \text{sigm}(\mathbf{c} + \mathbf{W}^* \mathbf{h}(x))\end{aligned}$$


For binary units

Encoder


$$\begin{aligned}h(x) &= g(a(x)) \\ &= \text{sigm}(\mathbf{b} + \mathbf{W}x)\end{aligned}$$

- In general it is hard to prove optimality results in nonlinear case.

Autoencoders (Example)



Autoencoder Example



- Need additional constraints to avoid learning an identity.
- In general, Encoder and Decoder filters can be different.

Loss Function (Examples)

- Reproduce (reconstruct) the input at the output layer
- $f(x)$ is the output of autoencoder (model prediction)
- **Loss function** for binary inputs

$$l(f(\mathbf{x})) = - \sum_k (x_k \log(\hat{x}_k) + (1 - x_k) \log(1 - \hat{x}_k))$$

- Cross-entropy error function (reconstruction loss) $f(\mathbf{x}) \equiv \hat{\mathbf{x}}$
- **Loss function** for real-valued inputs

$$l(f(\mathbf{x})) = \frac{1}{2} \sum_k (\hat{x}_k - x_k)^2$$

- sum of squared differences (reconstruction loss)
- we use a linear activation function at the output

Loss Function

- Parameter gradients are obtained by back-propagating the gradient $\nabla_{\hat{\mathbf{a}}(\mathbf{x}^{(t)})} l(f(\mathbf{x}^{(t)}))$ like in a regular network
 - Important: when using tied weights ($\mathbf{W}^* = \mathbf{W}^\top$), $\nabla_{\mathbf{W}} l(f(\mathbf{x}^{(t)}))$ is the sum of two gradients (Show this as an exercise).
 - this is because \mathbf{W} is present in the encoder and in the decoder

Autoencoder example

Encoder architecture:

Convolutional layer

16 output channels , kernel size = 3, stride=3, padding=1

ReLU

MaxPooling

size = 2X2, stride=2

Convolutional layer

8 output channels, kernel size =3, stride=2, padding=1

ReLU

MaxPooling

size = 2, stride=1

Decoder architecture:

Convolutional Transpose layer

8 input channels , 16 output channels, kernel size=3, stride=2, padding=1

ReLU

Convolutional Transpose layer

16 input chennels, 8 output channels, kernel size =5, stride=3, padding=1

ReLU

Convolutional Transpose layer

8 input channels , 1 output channel, kernel size=2, stride=2, padding=1

Tanh

Example: MNIST



100 samples from test
data set



100 generated samples using
Autoencoder

- 20 epochs, batch size = 128, Adam optimizer, learning rate = 0.001
- Normalized images with mean 0.5 and std equals to 0.5

Autoencoder

- How to extract meaningful hidden features?
 1. Undercomplete Representation
 2. Injecting noise to the input (Denoising Autoencoder)
 3. Penalizing the solution (Contractive Autoencoders)



Undercomplete Representation

- Hidden layer is undercomplete if smaller than the input layer (bottleneck layer, e.g. dimensionality reduction):

- hidden layer “compresses” the input
- may compress well only for the training distribution


- Hidden units will be

- good features for the training distribution
- may not be robust to other types of input



Overcomplete Representation

- Hidden layer is **overcomplete** if greater than the input layer
 - no compression in hidden layer
 - In some cases each hidden unit may copy a different input component
- No guarantee that the hidden units will extract **meaningful structure**




Denoising Autoencoder

- Idea: representation should be robust to introduction of noise:


- random assignment of subset of inputs to 0, with probability ν
- Similar to dropouts on the input layer
- Similar idea for Gaussian additive noise

- Reconstruction $\hat{\mathbf{x}}$ computed from the corrupted input $\tilde{\mathbf{x}}$

- Loss function compares $\hat{\mathbf{x}}$ reconstruction with the noiseless input \mathbf{x}



Denoising Autoencoder




Contractive Autoencoder

- Alternative approach to avoid **uninteresting solutions**
 - add an **explicit term** in the loss that penalizes that solution

You want denoising behavior without explicitly corrupting inputs

- We wish to extract features that only reflect variations observed in the training set

- we'd like to be invariant to the other variations



Contractive Autoencoder

- Consider the following loss function:

$$l(f(\mathbf{x}^{(t)})) + \lambda \underbrace{||\nabla_{\mathbf{x}^{(t)}} \mathbf{h}(\mathbf{x}^{(t)})||_F^2}_{\text{Jacobian of Encoder}}$$

$\underbrace{l(f(\mathbf{x}^{(t)}))}_{\text{Reconstruction Loss}}$

- Example for binary observations:


$$l(f(\mathbf{x}^{(t)})) = - \sum_k \left(x_k^{(t)} \log(\hat{x}_k^{(t)}) + (1 - x_k^{(t)}) \log(1 - \hat{x}_k^{(t)}) \right)$$

$$||\nabla_{\mathbf{x}^{(t)}} \mathbf{h}(\mathbf{x}^{(t)})||_F^2 = \sum_j \sum_k \left(\frac{\partial h(\mathbf{x}^{(t)})_j}{\partial x_k^{(t)}} \right)^2$$

Encoder throws
away all information


Autoencoder attempts to
preserve all information

Autoencoder



- Details of what goes inside the encoder and decoder matter!


Predictive Sparse Decomposition



At training time

$$\min_{D, W, z} \underbrace{\|Dz - x\|_2^2 + \lambda|z|_1}_{\text{Decoder}} + \underbrace{\|\sigma(Wx) - z\|_2^2}_{\text{Encoder}}$$

Stacked Autoencoders



Stacked Autoencoders

