

Variational Autoencoders

Vahid Tarokh


ECE 685D, Fall 2025

Generative Models-Review

- We discussed about two learning paradigms : generative and discriminative
- In generative paradigm the learner tries to learn a joint distribution over all the variables.
- A generative model simulates how the data is generated in the real world.
- Why do we need generative model:
 - Encoding the laws of physics and other constraints into the generative process
 - Expresses causal relations of the world, so we can generalize better to new situations than mere correlations.
 - Building useful abstractions of the world
 - Useful for unsupervised representation learning

VAE-Overview

- Variational Autoencoders:
 - Allow us to design complex generative models of data and fit them to large datasets.
 - Can generate fake data/images (e.g., fake celebrity faces) and fake high-resolution digital artwork.




VAE-Overview

- These models also yield state-of-the-art machine learning results in image generation.
- VAEs bridges graphical models and the deep learning
- The general construction is based on generative probability models and **does not really need to be implemented using Neural Networks.**
- Variational Autoencoders (VAEs) based on Deep Learning were proposed in 2013.
- But since this is a course on Neural Networks, we will mostly implement them using Neural Networks.

Ideal Generative Models

- The fundamental problem with autoencoders:
 - latent encoded set where their encoded vectors lie, may not be contiguous, or allow easy interpolation.




- Training an autoencoder on the MNIST dataset
- Visualizing the encodings from a 2D latent space
- The formation of distinct clusters.

<https://towardsdatascience.com/intuitively-understanding-variational-autoencoders-1bfe67eb5daf>


Optimizing only based on the maximum likelihood (reconstruction loss)

Ideal Generative Models

- What we may ideally want: the encodings are close to being “contiguous” (intuitively speaking) while still being distinct,
- This allows smooth interpolation and enables construction of *new* samples.



Desired encoding



Undesired encoding

Parameterization of conditional distributions with Neural Networks

- Probabilistic models based on neural networks are computationally scalable
 - Using stochastic gradient-based optimization
 - Scaling to large models and large datasets
- An example:
 - Parameterizing a categorical distribution (cat), $P_{\theta}(y|x)$ using neural networks (NN)
 - y : class label, conditioned on
 - x : input image
 - θ : the parameter vector of the distribution



$NN(x)$

$$\begin{aligned} p(y = \text{dog}|x) &= 0.979 \\ p(y = \text{cat}|x) &= 0.02 \\ p(y = \text{deer}|x) &= 0.001 \\ &\vdots \end{aligned}$$

Deep Latent-Variable Model (DLVM)

- DLVM denote a latent variable model $p_{\theta}(x, z)$ where distributions are parameterized by neural networks.
- Advantage of DLVM:
 - Even having simple factors (prior or conditional distribution) such as conditional Gaussian, the **marginal likelihood**, $p_{\theta}(x)$ can be arbitrarily complex

DLVM is attractive for approximating complicated underlying distributions

Example: DLVM for multivariate Bernoulli data

- An example from (Kingma and Welling, 2014) for binary data \mathbf{x} and with a spherical Gaussian latent space,

$$p(\mathbf{z}) = \mathcal{N}(\mathbf{z}; 0, \mathbf{I}) \quad \text{Latent variable distribution}$$

$$\mathbf{p} = \text{DecoderNeuralNet}_{\theta}(\mathbf{z}) \quad \text{Parametrized Bernoulli with NN}$$

$$\log p(\mathbf{x}|\mathbf{z}) = \sum_{j=1}^D \log p(x_j|\mathbf{z}) = \sum_{j=1}^D \log \text{Bernoulli}(x_j; p_j)$$

$$= \sum_{j=1}^D x_j \log p_j + (1 - x_j) \log(1 - p_j)$$

D denotes the dimension of \mathbf{x} and $0 \leq p_j$'s ≤ 1 .

Disadvantage of DLVM

- Intractability of Marginal likelihood (evidence)
$$\int p_{\theta}(x | z) p(z) dz$$

➤ So, we cannot differentiate it w.r.t. its parameters and optimize it
- This means the intractability of the posterior, $p(z|x)$
- Solution: Estimate posterior distribution using Variational Inference (variational Bayes)

VAE Framework

- Using Variational Autoencoders (VAEs), we can efficiently optimize DLVMs jointly with a corresponding inference model using SGD.
- In the probability model framework, a **variational auto-encoder** (VAE) is assumed to be modeled by a **specific probability model** of data \mathbf{x} and latent variables \mathbf{z} .
- We can write the joint probability of the model as $p_{\theta}(\mathbf{x}, \mathbf{z}) = p_{\theta}(\mathbf{x} | \mathbf{z}) p(\mathbf{z})$
- The **generative process** can be written as follows:
 - For each data-point i :
 - Draw latent variables $\mathbf{z}_i \sim p(\mathbf{z})$
 - Draw data-point $\mathbf{x}_i \sim p_{\theta}(\mathbf{x} | \mathbf{z})$

The Goal

- To get around of directly computing the intractable posterior, one can introduce parametric *inference model* $q_\lambda(z|x)$ to approximate the true posterior:

$$q_\lambda(z|x) \approx p_\theta(z|x)$$

- This model is also called an *encoder* or .
- The parameters of this inference model, λ also called the *variational parameters*.
- Using the above approximation, we can optimize the marginal likelihood.
- For instance if $q_\lambda(z|x)$ are Gaussian, the λ is the mean vector and co-variance matrix.

VAE Framework




Image source: Kingma and Welling, 2019

Variational Parameters

- Parameterizing the distribution of approximate posterior, $q_\lambda(z|x)$ by a deep neural network encoder:

➤ The variational parameters include the weights and biases of the neural network:

$$(\mu, \log \sigma) = EncoderNN_\phi(x)$$

$$q_\lambda(z|x) = \mathcal{N}(\mathbf{z}; \mu, diag(\sigma))$$

➤ Using a single encoder neural network to perform posterior inference over all datapoints (shared parameters).

- This is called *amortized variational inference*

➤ In traditional approximate inference models, variational parameters are not shared

Question: How to choose λ ?

- The optimization objective is reverse KL:
 - The *evidence lower bound*, abbreviated as ELBO (*variational lower bound*)

$$\begin{aligned} & D(q_\lambda(z|x) || p_\theta(z|x)) \\ &= E_{q_\lambda(z|x)}[\log q_\lambda(z|x)] + \log(p_\theta(x)) \\ &\quad - E_{q_\lambda(z|x)}[\log p_\theta(x, z)] \end{aligned}$$

- Hence:

$$\begin{aligned} \log(p(x)) &= D(q_\lambda(z|x) || p_\theta(z|x)) + \\ & E_{q_\lambda(z|x)}[\log p_\theta(x, z)] - \\ & E_{q_\lambda(z|x)}[\log q_\lambda(z|x)] \end{aligned}$$

Evidence Lower Bound (ELBO)

- This means that minimizing the KL-distance is equivalent to maximizing

$$ELBO = E_{q_\lambda(z|x)}[\log p_\theta(x,z)] - E_{q_\lambda(z|x)}[\log q_\lambda(z|x)]$$

- Assume that no two data points share their latent variables with each other then ELBO decomposes into the sum of

$$ELBO_i = E_{q_\lambda(z|x_i)}[\log p_\theta(x,z)] - E_{q_\lambda(z|x_i)}[\log q_\lambda(z|x_i)]$$

The Main Message

- This is equal to


$$ELBO_i = E_{q_\lambda(z|x_i)} [\log p_\theta(x_i | z)] - D(q_\lambda(z|x_i) || p(z))$$

Likelihood term (reconstruction part)

Closeness of encoding to $p(z)$
(typically Gaussian)

- Typically $p(z)$ is selected to be standard normal distribution. Then if we have a parametric model class for $p_\theta(x_i | z)$, we can maximize the objective function $\sum_i ELBO_i$ over parameters θ and λ using Stochastic Gradient Ascent.
- This process is true regardless of if model classes $p_\theta(x_i | z)$ and $q_\lambda(z|x)$ are given by deep Neural Networks or not!

Computational Flow In a Variational Autoencoder



Generation of New Data

- Let us introduce some names:
 - $q_\lambda(z|x)$ is referred to as the **encoder**
 - $p_\theta(x|z)$ is referred to as the **decoder**
- If we have an encoder and decoder designed (optimized as before) and have calculated all the optimized parameters, then to make a new value of x (generate data)
 - Generate z according to $p(z)$ (Typically standard Gaussian)
 - Generate x according to $p_\theta(x|z)$


Connecting to Neural Networks

- But this course is about Neural Networks.
- Idea: Use DNNs somehow for encoders and Decoders.
 - Encoding Neural Network: Up on observing x , the neural network outputs parameters λ
 - Decoding Neural Network: Up on observing z , the neural network outputs parameters θ
- Equivalently we will have to learn the weights λ and θ of encoding and decoding DNNs using SGD to minimize $-ELBO$.
- We can present this in a fancier way if we wish.

VAEs from Multilayer DNNs

Deep VAE


- Consider multilevel features:



Generative (Decoder) Network

$$p(\mathbf{x}|\boldsymbol{\theta}) = \sum_{\mathbf{h}^1, \dots, \mathbf{h}^L} p(\mathbf{h}^L|\boldsymbol{\theta})p(\mathbf{h}^{L-1}|\mathbf{h}^L, \boldsymbol{\theta}) \cdots p(\mathbf{x}|\mathbf{h}^1, \boldsymbol{\theta})$$

Each term may denote a complicated nonlinear relationship




- $\boldsymbol{\theta}$ denotes parameters of decoder.
- L is the number of **stochastic** layers.
- Sampling and probability evaluation is tractable for each $p(\mathbf{h}^\ell|\mathbf{h}^{\ell+1})$.

Recognition (Encoder) Network

- The recognition model (**encoder**) is defined in terms of an analogous factorization:

$$q(\mathbf{h}|\mathbf{x}, \lambda) = q(\mathbf{h}^1|\mathbf{x}, \lambda) q(\mathbf{h}^2|\mathbf{h}^1, \lambda) \dots q(\mathbf{h}^L|\mathbf{h}^{L-1}, \lambda)$$



Each term may denote a complicated nonlinear relationship

We assume that

$$\mathbf{h}^L \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$$

The conditionals:

$$\begin{aligned} p(\mathbf{h}^\ell | \mathbf{h}^{\ell+1}) \\ q(\mathbf{h}^\ell | \mathbf{h}^{\ell-1}) \end{aligned}$$

Using SGD for training Both Networks

- We can train decoder and encoder networks using SGD algorithm for minimizing the $\mathcal{L}_{\theta,\lambda}(\mathbf{x}) \triangleq -\text{ELBO}$ (\mathcal{D} denotes a mini-batch):

$$\mathcal{L}_{\theta,\lambda}(\mathbf{x}) = \sum_{\mathbf{x} \in \mathcal{D}} \mathcal{L}_{\theta,\lambda}(\mathbf{x})$$

$$\mathcal{L}_{\theta,\lambda}(\mathbf{x}) = E_{q_\lambda(z|x)}[\log p_\theta(x,z)] - E_{q_\lambda(z|x)}[\log q_\lambda(z|x)]$$

- Obtaining Unbiased gradient w.r.t the decoder parameters, θ is simple:

$$\begin{aligned}\nabla_\theta(\mathcal{L}_{\theta,\lambda}(\mathbf{x})) &= \nabla_\theta(E_{q_\lambda(z|x)}[\log p_\theta(x,z)] - E_{q_\lambda(z|x)}[\log q_\lambda(z|x)]) \\ &= E_{q_\lambda(z|x)}[\nabla_\theta(\log p_\theta(x,z) - \log q_\lambda(z|x))] \\ &\simeq \nabla_\theta(\log p_\theta(x,z) - \log q_\lambda(z|x)) \\ &= \nabla_\theta(\log p_\theta(x,z))\end{aligned}$$

Using SGD for training Both Networks

- Unbiased gradient w.r.t the variational parameters, λ is more difficult:
 - The Expectation is taken w.r.t the $q_\lambda(z|x)$, which is a function of λ
!!!

$$\begin{aligned}\nabla_\lambda(\mathcal{L}_{\theta,\lambda}(x)) &= \nabla_\lambda(E_{q_\lambda(z|x)}[\log p_\theta(x,z)] - E_{q_\lambda(z|x)}[\log q_\lambda(z|x)]) \\ &\neq E_{q_\lambda(z|x)}[\nabla_\lambda(\log p_\theta(x,z) - \log q_\lambda(z|x))]\end{aligned}$$

- Fortunately, for the continuous r.v's, the unbiased estimator of the gradient can be obtained through the reparameterization trick.

Reparameterization Trick

- This trick is essentially the application of change of variables:
 - Let z be a r.v distributed as $q_\lambda(z|x)$
 - Assume, there is a differentiable and invertible function h such that

$$z = h(\epsilon, x, \lambda)$$


- Where $\epsilon \sim p(\epsilon)$ is another random variable which is independent of x and λ . Hence:

$$E_{q(z|x)}[f(z)] = E_{p(\epsilon)}[f(z)]$$

- Now, we can safely exchange the gradient and expectation ($f(z) = \log q_\lambda(z|x)$):

$$\begin{aligned}\nabla_\lambda E_{q_\lambda(z|x)}[\log q_\lambda(z|x)] &= E_{p(\epsilon)} [\nabla_\lambda (\log q_\lambda(z|x))] \\ &= \nabla_\lambda E_{p(\epsilon)}[\log q_\lambda(z|x)] \simeq \nabla_\lambda (\log q_\lambda(z|x))\end{aligned}$$

Illustration of Reparameterization Trick



- The reparametrized ELBO estimator is referred to as the **Stochastic Gradient Variational Bayes (SGVB)** estimator (Kingma and Welling, 2014)

Computing the Distribution of The Approximate Posterior After Change of Variables

- The goal is to compute the distribution of $\log q_\lambda(z|x)$ by changing z as a function of ϵ .
- This is easy task as long as the function h is chosen appropriately
- Using the change of variable rule in probability:

$$\log q_\lambda(z|x) = \log p(\epsilon) - \log |\det\left(\frac{\partial z}{\partial \epsilon}\right)|$$

- $\frac{\partial z}{\partial \epsilon}$ denotes the Jacobian matrix which is computed through $\mathbf{z} = h(\epsilon, x, \lambda)$.

Two common and simple choices for function h

1. Factorized Gaussian posterior:

- We can assume that $q_\lambda(z|x)$ is given by

$$q_\lambda(z|x) \sim \mathcal{N}(\mu, \text{diag}(\sigma^2))$$

- In this case, variational parameters, and the approximate posterior, $q_\lambda(z|x)$ are given by:

$$(\mu, \log \sigma) = \text{EncoderNN}_\lambda(x)$$

$$q_\lambda(z|x) = \prod_i \mathcal{N}(z_i; \mu_i, \sigma_i^2) \quad \text{Univariate Gaussian Distribution}$$

- We take h as an affine function of ϵ as follows:

$$\epsilon \sim \mathcal{N}(0, I)$$

$$z = \mu + \sigma \odot \epsilon$$

$$\text{Jacobian: } \frac{\partial z}{\partial \epsilon} = \text{diag}(\sigma) \xrightarrow{\text{yields}} \log \left| \det \left(\frac{\partial z}{\partial \epsilon} \right) \right| = \sum_i \log(\sigma_i)$$

Two common and simple choices for function h

2. Full-covariance Gaussian posterior

- We can assume that $q_\lambda(z|x)$ is given by

$$q_\lambda(z|x) \sim \mathcal{N}(\mu, \Sigma)$$
$$\Sigma = LL^T \quad \text{Cholesky Decomposition}$$

- L is a lower or upper triangle matrix with non-zero entries in the diagonal.

- In this case, variational parameters is given by:

$$(\mu, \log \sigma, L') = \text{EncoderNN}_\lambda(x)$$

$$L = L_{mask} \odot L' + \text{diag}(\sigma)$$

- L_{mask} is a 0-1 matrix with zero's on and above (lower) diagonal.

- We take h as an affine function of ϵ as follows:

$$\epsilon \sim \mathcal{N}(0, I)$$

$$z = \mu + L\epsilon$$

Jacobian: $\frac{\partial z}{\partial \epsilon} = L \xrightarrow{\text{yields}} \log \left| \det \left(\frac{\partial z}{\partial \epsilon} \right) \right| = \sum_i \log |L_{ii}|$

Computing the Gradients

- The gradient w.r.t the parameters: both recognition and generative:

$$\nabla_{\theta, \lambda} \mathbb{E}_{z \sim q(z|x, \lambda)} \left[\log \frac{p(x, z | \theta)}{q(z | x, \lambda)} \right]$$

$$= \nabla_{\theta, \lambda} \mathbb{E}_{\epsilon^1, \dots, \epsilon^L \sim \mathcal{N}(\mathbf{0}, I)} \left[\log \frac{p(x, h(\epsilon, x, \theta) | \theta)}{q(h(\epsilon, x, \lambda) | x, \lambda)} \right]$$

$$= \mathbb{E}_{\epsilon^1, \dots, \epsilon^L \sim \mathcal{N}(\mathbf{0}, I)} \left[\nabla_{\theta, \lambda} \log \frac{p(x, h(\epsilon, x, \theta) | \theta)}{q(h(\epsilon, x, \lambda) | x, \lambda)} \right]$$



Gradients can be
computed by backprop

The mapping h is a deterministic
neural net for fixed ϵ .

Computing the Gradients

- The gradient w.r.t the parameters: recognition and generative:

$$\nabla_{\theta, \lambda} \mathbb{E}_{\mathbf{z} \sim q(\mathbf{z} | \mathbf{x}, \lambda)} \left[\log \frac{p(\mathbf{x}, \mathbf{z} | \boldsymbol{\theta})}{q(\mathbf{z} | \mathbf{x}, \lambda)} \right] = \mathbb{E}_{\boldsymbol{\epsilon}^1, \dots, \boldsymbol{\epsilon}^L \sim \mathcal{N}(\mathbf{0}, \mathbf{I})} \left[\nabla_{\theta, \lambda} \log \frac{p(\mathbf{x}, g(\boldsymbol{\epsilon}, \mathbf{x}, \boldsymbol{\theta}) | \boldsymbol{\theta})}{q(h(\boldsymbol{\epsilon}, \mathbf{x}, \lambda) | \mathbf{x}, \lambda)} \right]$$

is h , not g

- Approximate expectation by generating k samples from $\boldsymbol{\epsilon}$:

$$\frac{1}{k} \sum_{i=1}^k \nabla_{\theta, \lambda} \log w(\mathbf{x}, \mathbf{h}(\boldsymbol{\epsilon}_i, \mathbf{x}, \lambda), \boldsymbol{\theta})$$

- Where we defined unnormalized **importance weights**:

$$w(\mathbf{x}, \mathbf{h}, \boldsymbol{\theta}) = p(\mathbf{x}, \mathbf{h} | \boldsymbol{\theta}) / q(\mathbf{h} | \mathbf{x}, \lambda)$$

Relation between ML and ELBO

- With i.i.d samples from dataset \mathcal{D} , the maximum likelihood principle is given by:

$$\log p_\theta(\mathcal{D}) = \frac{1}{N} \sum_{\mathbf{x} \in \mathcal{D}} \log p_\theta(\mathbf{x}) = E_{q_{\mathcal{D}}(\mathbf{x})} \log p_\theta(\mathbf{x})$$

- $q_{\mathcal{D}}(\mathbf{x})$ is the empirical distribution of data
- Recall that

$$\text{Max}_\theta \log p_\theta(\mathcal{D}) = \text{Min}_\theta D_{kl}(q_{\mathcal{D}}(\mathbf{x}) || p_\theta(\mathbf{x}))$$



- Combining the empirical distribution of data and the inference model in VAE, $q_{\lambda, \mathcal{D}}(\mathbf{x}, \mathbf{z}) = q_{\mathcal{D}}(\mathbf{x})q_\lambda(\mathbf{z}|\mathbf{x})$, one can show that

$$D_{kl}(q_{\lambda, \mathcal{D}}(\mathbf{x}, \mathbf{z}) || p_\theta(\mathbf{x}, \mathbf{z})) \geq D_{kl}(q_{\mathcal{D}}(\mathbf{x}) || p_\theta(\mathbf{x}))$$

- ELBO can be thought as a maximum likelihood objective *in an augmented space*, (\mathbf{x}, \mathbf{z})

Trade-off in the VAE loss

$$E_{q_\lambda(z|x_i)} [\log p(x_i | z)] - D(q_\lambda(z|x) || p(z))$$




Optimizing using the second term (KL divergence) in the loss

<https://towardsdatascience.com/intuitively-understanding-variational-autoencoders-1bfe67eb5daf>

Optimizing using both reconstruction loss (likelihood term) and KL divergence term in the loss

VAE Architectures for MNIST



MNIST Experiment

Reconstruction



1st epoch



5th epoch



10th epoch

- Adam optimizer, learning rate=0.001, batch size = 128, 10 epochs, no image normalization
- Reconstruction of images in the output of VAE in different epochs

MNIST Experiment

Sampling (generating)



1st epoch



5th epoch



10th epoch

- Generating of images in the output of VAE in different epochs by sampling a fixed random vector $Z \in \mathbb{R}^{128 \times 20}$ with $Z_i \sim \mathcal{N}(0, I_{20 \times 20})$, $i = 1, 2, \dots, 128$


Other Variants

Importance Weighted Autoencoders

- Consider the following k -sample importance weighting of the log-likelihood ($z_i = h_i = h(\epsilon_i, \mathbf{x}, \lambda)$):

$$\mathcal{L}_k(\mathbf{x}) = \mathbb{E}_{\mathbf{h}_1, \dots, \mathbf{h}_k \sim q(\mathbf{h}|\mathbf{x})} \left[\log \frac{1}{k} \sum_{i=1}^k \frac{p(\mathbf{x}, \mathbf{h}_i)}{q(\mathbf{h}_i|\mathbf{x})} \right]$$

$$= \mathbb{E}_{\mathbf{h}_1, \dots, \mathbf{h}_k \sim q(\mathbf{h}|\mathbf{x})} \left[\log \frac{1}{k} \sum_{i=1}^k w_i \right]$$



Unnormalized
importance weights

where $\mathbf{h}_1, \dots, \mathbf{h}_k$ are sampled
from the recognition network.

Importance Weighted Autoencoders

- Consider the following k -sample importance weighting of the log-likelihood ($z_i = h_i = h(\epsilon_i, \mathbf{x}, \lambda)$):

$$\mathcal{L}_k(\mathbf{x}) = \mathbb{E}_{\mathbf{h}_1, \dots, \mathbf{h}_k \sim q(\mathbf{h}|\mathbf{x})} \left[\log \frac{1}{k} \sum_{i=1}^k \frac{p(\mathbf{x}, \mathbf{h}_i)}{q(\mathbf{h}_i|\mathbf{x})} \right]$$

- This is a lower bound on the marginal log-likelihood:

$$\mathcal{L}_k(\mathbf{x}) = \mathbb{E} \left[\log \frac{1}{k} \sum_{i=1}^k w_i \right] \leq \log \mathbb{E} \left[\frac{1}{k} \sum_{i=1}^k w_i \right] \stackrel{\text{Calculate this}}{=} \log p(\mathbf{x})$$

- Special Case of $k=1$:** Same as standard VAE objective.
- Using more samples → Improves the tightness of the bound.

Tighter Lower Bound

- Using more samples can only improve the tightness of the bound.
- For all k , the lower bounds satisfy:

$$\log p(\mathbf{x}) \geq \mathcal{L}_{k+1}(\mathbf{x}) \geq \mathcal{L}_k(\mathbf{x})$$

- Moreover if $p(\mathbf{h}, \mathbf{x})/q(\mathbf{h}|\mathbf{x})$ is bounded, then:

$$\mathcal{L}_k(\mathbf{x}) \rightarrow \log p(\mathbf{x}), \quad \text{as } k \rightarrow \infty$$

Computing the Gradients

- We can use the unbiased estimate of the gradient using reparameterization trick:

$$\begin{aligned}\nabla_{\theta, \lambda} \mathcal{L}_k(\mathbf{x}) &= \nabla_{\theta, \lambda} \mathbb{E}_{\mathbf{h}_1, \dots, \mathbf{h}_k \sim q(\mathbf{h} | \mathbf{x})} \left[\log \frac{1}{k} \sum_{i=1}^k w_i \right] \\ &= \mathbb{E}_{\boldsymbol{\epsilon}_1, \dots, \boldsymbol{\epsilon}_k} \left[\nabla_{\theta, \lambda} \log \frac{1}{k} \sum_{i=1}^k w(\mathbf{x}, h(\boldsymbol{\epsilon}_i, \mathbf{x}, \lambda), \boldsymbol{\theta}) \right] \\ &= \mathbb{E}_{\boldsymbol{\epsilon}_1, \dots, \boldsymbol{\epsilon}_k} \left[\sum_{i=1}^k \tilde{w}_i \nabla_{\theta, \lambda} \log w(\mathbf{x}, h(\boldsymbol{\epsilon}_i, \mathbf{x}, \lambda), \boldsymbol{\theta}) \right]\end{aligned}$$

- Where we define normalized importance weights:

$$\tilde{w}_i = w_i / \sum_{i=1}^k w_i, \quad \text{where } w_i = \frac{p(\mathbf{x}, \mathbf{h}_i)}{q(\mathbf{h}_i | \mathbf{x})}$$

IWAEs vs. VAEs

- Draw k -samples from the recognition network $q_\lambda(\mathbf{h}|\mathbf{x})$
 - or k -sets of auxiliary variables ϵ .
- Obtain the following Monte Carlo estimate of the gradient (IWAE):

$$\nabla_{\theta,\lambda} \mathcal{L}_k(\mathbf{x}) \approx \sum_{i=1}^k \tilde{w}_i [\nabla_{\theta,\lambda} \log w(\mathbf{x}, \mathbf{h}(\epsilon_i, \mathbf{x}, \lambda), \theta)]$$

- Compare this to the VAE's estimate of the gradient (VAE):

$$\nabla_{\theta,\lambda} \mathcal{L}(\mathbf{x}) \approx \frac{1}{k} \sum_{i=1}^k [\nabla_{\theta,\lambda} \log w(\mathbf{x}, \mathbf{h}(\epsilon_i, \mathbf{x}, \lambda), \theta)]$$

Conditional VAE

Conditional VAE (CWAE)

- No control on the data generation process on VAE
 - e.g., We want to generate only a digit 2
- Conditioning all the distributions on what we want, the objective of CWAE (conditional ELBO):

$$E_{q_\lambda(z|x)}[\log p_\theta(x|z, c) - D_{kl}(q_\lambda(z|x, c) || p(z|c))]$$

- c denotes the conditioning vector (e.g., a code for digit 2).

CWAE in Practice

- Conditioned image generation




Image source: Yeh et al., lecture notes, 2017

β -VAE

Disentangled Latent Features— β -VAE

- A latent variable z is called disentangled factor if it is
 - Only sensitive to one single generative factor
 - Relatively invariant to other factors
- Advantage:
 - Interpretability
 - Easy generalization to a different task
- Example: Disentangled factors in Human face
 - Skin color
 - Hair length
 - Having glass or not

Disentangled Latent Features— β -VAE

- β -VAE was proposed by Higgins et al., 2017 to discover the disentangled latent representation:

$$\begin{aligned} & \underset{\theta, \lambda}{\operatorname{Max}} E_{q_\lambda(z|x)} [\log p_\theta(x, z)] \\ & \text{s.t. } D_{kl}(q_\lambda(z|x) || p(z)) < \delta \end{aligned}$$

- If $\beta=1$, this is regular VAE
- If $\beta > 1$, a stronger constraint on the latent factors; hence, limiting the representation capacity of z
- The larger β , the better discovering of disentangled latent features **Using Lagrangian form**




$$\underset{\theta, \lambda}{\operatorname{Max}} E_{q_\lambda(z|x)} [\log p_\theta(x, z) - \beta D_{kl}(q_\lambda(z|x) || p(z))]$$

β -VAE in Practice


- β -VAE: $\beta = 250$
- VAE: $\beta = 1$

- Disentangled latent factors: Azimuth, Emotion, Hair
- β -VAE learns to **disentangle** factors
- VAE learns an **entangled** representation
 - Entangling azimuth with emotion
 - presence of glasses and gender



β -VAE in Practice


- β -VAE: $\beta = 5$
- VAE: $\beta = 1$
- Disentangled latent factors: Azimuth, Chair Width, Leg Style
- β -VAE learns to **disentangle** factors
- VAE learns an **entangled** representation
 - Entangling chair width with azimuth and leg style



Vector Quantized-VAE

Vector Quantized-Variational AutoEncoder (VQ-VAE)

- Learning a discrete latent variable by the encoder



- Mapping K -dimensional vectors into a finite set of “code” vectors.
- Similar to K-means algorithm

Image source: van den Oord, et al. 2017