

**Midterm Exam I**  
**ECE 685D– Introduction to Deep Learning**  
**Fall 2024**

Instructor: Prof. Vahid Tarokh  
ECE Department, Duke University

Oct 9<sup>th</sup> 2024  
10:05 AM - 11:20 AM

Name: \_\_\_\_\_  
Duke ID: \_\_\_\_\_

---

This exam contains 10 pages and 10 questions. This is a closed-book exam. No exam aids are permitted except for a one-sided letter-sized cheat sheet. Communication with others is strictly prohibited.

**Distribution of Marks**

Question	Points	Score
1	3	
2	3	
3	3	
4	3	
5	3	
6	3	
7	5	
8	25	
9	35	
10	17	
Total:	100	

**Definition of Leaky ReLU function.** Given the non-negative coefficient  $\alpha$  (i.e.,  $\alpha < 1$ ) and the input  $x$ , the leaky ReLU is defined as:

$$\text{LeakyReLU}(x, \alpha) = \begin{cases} x & \text{if } x \geq 0 \\ \alpha x & \text{if } x < 0 \end{cases}$$

For the multiple-choice questions, circle ALL of the correct answers. These questions can have more than one correct answer.

1. (3 points) After training a convolutional neural network, you observe a large gap between the training accuracy (e.g., 99%) and the test accuracy (e.g., 50%). Which of the following methods is commonly used to resolve this issue?

- A. (a) Dropout  
 (b) Leaky ReLU  
 (c) Sigmoid activation

2. (3 points) Which of the following activation functions can lead to vanishing gradients?

- A. C. (a) Sigmoid  
 (b) Leaky ReLU  
 (c) Tanh

3. (3 points) Which of the following is true about Batchnorm

- B. (a) Batchnorm is another way of performing dropout.  
 (b) Batchnorm makes the training process converge faster.  
 (c) Batchnorm is a non-linear transformation to center the dataset around the origin.

C. 4. (3 points) If the size of the input is  $64 \times 64 \times 16$  (i.e., number of channels = 16), how many parameters are there, including bias, in a convolution filter with kernel size=1, stride=1, padding=0?

- (a) 1  
 (b) 2  
 (c) 17

Weight param =  $K \times K \times \# \text{ channels}$ .

Total = weight param + # bias,  $l = \# \text{ filter}$ .

Each filter has its own bias

(d) 4097

*The # filters determine # output channels.*

5. (3 points) Consider an input of dimensions  $(n_h, n_w, n_c)$ , where  $n_c$  is the number of channels. A convolutional layer with kernel size=1, stride=1, padding=0 is applied to this input. After the convolution, a standard max-pooling layer (e.g., stride=2) is applied. Which of the following statements about the output after both operations is correct?

~~C.~~  
B.C

- (a) The convolution can help reduce the dimension of  $n_h, n_w$ , but not  $n_c$ .
- ~~(b)~~ The convolution can help reduce the dimension of  $n_c$ , but not  $n_h, n_w$ .
- (c) The standard maxpooling can help reduce the dimension of  $n_h, n_w$ , but not  $n_c$ .
- (d) The standard maxpooling can help reduce the dimension of  $n_c$ , but not  $n_h, n_w$ .

6. (3 points) Recall that the  $L_p$  norm of  $\mathbf{w} = [w_1, w_2, \dots, w_n]^T$  is defined as:

~~B.~~

$$\|\mathbf{w}\|_p = \left( \sum_i^n |w_i|^p \right)^{1/p}.$$

Which of the following regularization metrics is convex and most likely leads to weight sparsity?

	Convex	Sparsity	
(a) $L_{1/2}$	$L_{1/2}$	$\times$	$\checkmark$
(b) $L_1$	$L_1$	$\checkmark$	$\checkmark$
(c) $L_2$	$L_2$	$\checkmark$	$\times$
(d) $L_\infty$	$L_\infty$	$\checkmark$	$\times$

*strong sparsity  
Lasso, drives weights to 0.  
weights rarely to 0.*

7. (5 points) What are the advantages of using convolutional layers instead of fully connected layers for visual tasks (e.g., image classification, object detection)? Explain your answer.

*Fully connected layers have too many params.  
not practical to train.*

8. Consider an image  $X$  with 2 channels. The image is represented as a  $4 \times 4 \times 2$  matrix where the last dimension correspond to the number of channel of this image. We denote a kernel  $W$  with dimension  $3 \times 3 \times 2$ . The values of the image pixels and weights of the kernel are given below.

$$X[:, :, 0] = \begin{bmatrix} 3 & 1 & 1 & 1 \\ 0 & 2 & 1 & 0 \\ 1 & 0 & 1 & 1 \\ 1 & 0 & 1 & 3 \end{bmatrix} \quad W[:, :, 0] = \begin{bmatrix} -1 & 1 & 1 \\ 0 & -1 & 0 \\ -1 & 1 & 1 \end{bmatrix}$$

$$X[:, :, 1] = \begin{bmatrix} 1 & 1 & 2 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 1 & 2 & 1 \\ 0 & 0 & 2 & 1 \end{bmatrix} \quad W[:, :, 1] = \begin{bmatrix} -1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & -1 \end{bmatrix}$$

The output of the convolutional filter is given as:

$$Y = \text{LeakyReLU}(X' * W - \mathbf{1}_{4 \times 4}, 0.1), \quad (1)$$

where  $Y$  is the output image,  $X'$  is the input image after applying zero-padding around the edges (i.e. each channel is converted to a  $6 \times 6$  matrix such that a row of zeros is added to the top and bottom and a column of zeros is added to the left and right.),  $X' * W$  is the convolution between  $X'$  and  $W$  with stride size=1, and  $\mathbf{1}_{4 \times 4}$  is a  $4 \times 4$  matrix with elements equal to one. Here,  $\alpha$  is the coefficient of the LeakyReLU function, where the leaky ReLU function is defined as:

$$\text{LeakyReLU}(x, \alpha) = \begin{cases} x & \text{if } x \geq 0, \\ \alpha x & \text{if } x < 0. \end{cases}$$

- (a) (15 points) Compute the output  $Y$  of the image  $X$ .
- (b) (5 points) Apply max pooling on non-overlapping  $2 \times 2$  sub-matrices of the output image and compute the output.
- (c) (5 points) Apply average pooling on non-overlapping  $2 \times 2$  sub-matrices of the output image and compute the output.



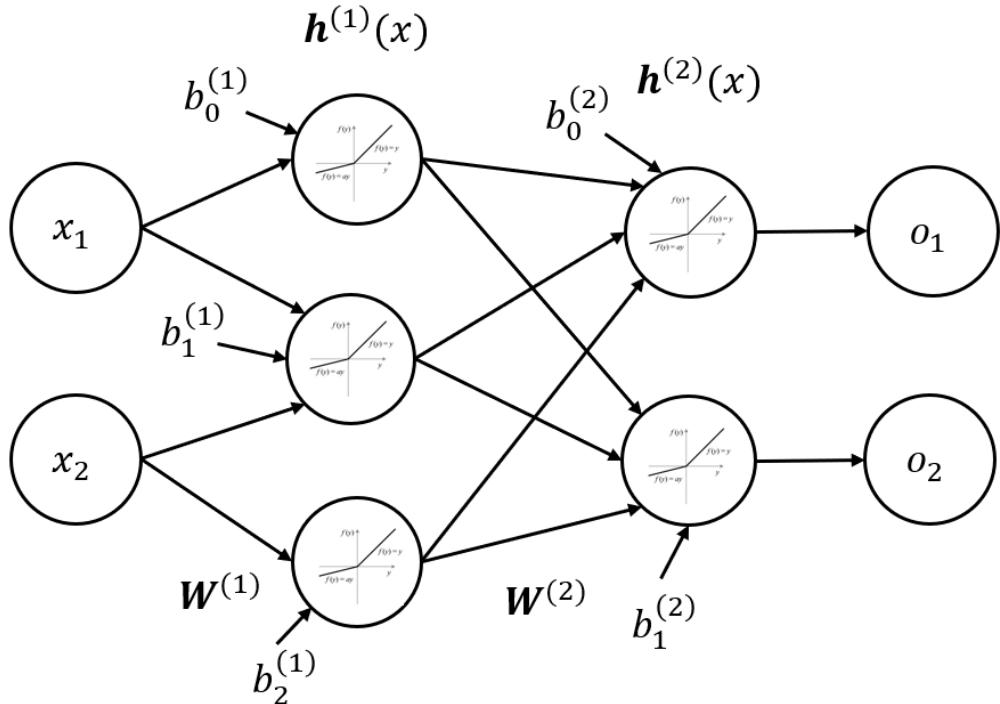
9. Given the neural network with  $L = 2$  hidden layers, the input units  $\mathbf{x} = [x_1, x_2]^T$ , and LeakyReLU activations (as shown below). The weights and bias of hidden units are denoted  $\mathbf{W}$  and  $\mathbf{b}$ . The hidden layers are defined as below:

$$\mathbf{h}^{(0)}(\mathbf{x}) = \mathbf{x} = [x_1, x_2]^T$$

$$\mathbf{h}^{(i)}(\mathbf{x}) = \sigma(\mathbf{b}^{(i)} + \mathbf{W}^{(i)}\mathbf{h}^{(i-1)}(\mathbf{x})), \text{ for } i = \{1, 2\}$$

where  $\sigma$  is the LeakyReLU activation with the coefficient  $\alpha = 0.5$ . The output  $\mathbf{o} = [o_1, o_2]^T$  is defined as:

$$\mathbf{o}(\mathbf{x}) = \mathbf{h}^{(2)}(\mathbf{x})$$



The values of the inputs are  $\mathbf{x} = [1, 2]^T$ . The ground truth values of the output are  $\mathbf{t} = [4, 1]^T$ . The weights of the network are given as follows:

$$\mathbf{W}^{(1)} = \begin{bmatrix} -0.5 & 0 \\ 0.5 & 1 \\ 0 & 1 \end{bmatrix} \quad \mathbf{W}^{(2)} = \begin{bmatrix} 0.5 & 2 & 0.5 \\ 1 & -0.5 & -2 \end{bmatrix}$$

$$\mathbf{b}^{(1)} = [2.5 \quad -0.5 \quad 1]^T \quad \mathbf{b}^{(2)} = [-5 \quad -1]^T$$

- (a) (10 points) Compute the output  $[o_1, o_2]^T$  of the input  $[x_1, x_2]^T$  using the network parameters as specified above. Write down all calculations of the intermediate layers.

- (b) (5 points) Compute the mean squared error (MSE) between the output  $[o_1, o_2]^T$  and the target ground truth  $[t_1, t_2]^T$ .
- (c) (5 points) Using the calculated MSE, update the weight  $W_{0,1}^{(2)}$  (i.e., entry  $\{0, 1\}$  of matrix  $\mathbf{W}^{(2)}$ ) using gradient descent and the backpropagation algorithm with the learning rate of 0.1. You do not need to simplify your answer.
- (d) (5 points) Similarly, update the weight  $W_{1,1}^{(2)}$  (i.e., entry  $\{1, 1\}$  of matrix  $\mathbf{W}^{(2)}$ ) using gradient descent and the backpropagation algorithm with the learning rate of 0.1. You do not need to simplify your answer.
- (e) (10 points) Update the weight  $W_{1,0}^{(1)}$  (i.e., entry  $\{1, 0\}$  of matrix  $\mathbf{W}^{(1)}$ ) using gradient descent and the backpropagation algorithm with a learning rate of 0.1. Write down the details of the backward pass. You do not need to simplify your answer.

$$\begin{aligned}
 (a1). \quad z^{(1)}(x) &= W^{(1)} h^{(0)}(x) + b^{(1)} = \begin{bmatrix} -0.5 & 0 \\ 0.5 & 1 \\ 0 & 1 \end{bmatrix} \cdot \begin{bmatrix} 1 \\ 2 \end{bmatrix} + \begin{bmatrix} 2.5 \\ -0.5 \\ 1 \end{bmatrix} \\
 &= \begin{bmatrix} -0.5 \\ 2.5 \\ 2 \end{bmatrix} + \begin{bmatrix} 2.5 \\ -0.5 \\ 1 \end{bmatrix} = \begin{bmatrix} 2 \\ 2 \\ 3 \end{bmatrix} \\
 h^{(1)}(x) &= \sigma(z^{(1)}(x)) = \begin{bmatrix} 2 \\ 2 \\ 3 \end{bmatrix} \\
 z^{(2)}(x) &= W^{(2)} \cdot h^{(1)}(x) + b^{(2)} = \begin{bmatrix} 0.5 & 2 & 0.5 \\ 1 & -0.5 & -2 \end{bmatrix} \cdot \begin{bmatrix} 2 \\ 2 \\ 3 \end{bmatrix} + \begin{bmatrix} -5 \\ -1 \end{bmatrix} \\
 &= \begin{bmatrix} 6.5 \\ -5 \end{bmatrix} + \begin{bmatrix} -5 \\ -1 \end{bmatrix} = \begin{bmatrix} 1.5 \\ -6 \end{bmatrix} \\
 h^{(2)}(x) &= \sigma(z^{(2)}(x)) = \begin{bmatrix} 1.5 \\ -3 \end{bmatrix} \\
 [o_1, o_2]^T &= h^2(x)^T = \begin{bmatrix} 1.5 & -3 \end{bmatrix}^T \\
 (b). \quad L &= \frac{1}{2} \sum_{i=1}^2 (o_i - t_i)^2 = \frac{1}{2} \cdot ((4 - 1.5)^2 + (1 + 3)^2) = 11.125.
 \end{aligned}$$

$$(c). \frac{\partial L}{\partial w_{0,1}^{(2)}} = \frac{\partial L}{\partial o_0} \cdot \frac{\partial o}{\partial h} \cdot \frac{\partial h}{\partial z} \cdot \frac{\partial z}{\partial w_{0,1}^{(2)}}$$

$$= (o_1 - t_1) \cdot 1 \cdot 1 \cdot 2 = -5.$$

$$w_{0,1}^{(2)} = w_{0,1}^{(2)} - 0.1 \times -5 = 2.5$$

$$(d). \frac{\partial L}{\partial w_{1,1}^{(2)}} = \frac{\partial L}{\partial o_1} \cdot \frac{\partial o}{\partial h} \cdot \frac{\partial h}{\partial z} \cdot \frac{\partial z}{\partial w_{1,1}^{(2)}}$$

$$= (o_2 - t_2) \cdot 1 \cdot 0.5 \cdot 2 = -4.$$

$$w_{1,1}^{(2)} = w_{1,1}^{(2)} - 0.1 \times -4 = -0.1$$

$$(e). \frac{\partial L}{\partial w_{1,0}^{(1)}} = \frac{\partial L}{\partial o_1} \cdot \frac{\partial o_1}{\partial h_1^{(2)}} \cdot \frac{\partial h_1^{(2)}}{\partial z_1^{(2)}} \cdot \frac{\partial z_1^{(2)}}{\partial h_1^{(1)}} \cdot \frac{\partial h_1^{(1)}}{\partial z_2^{(1)}} \cdot \frac{\partial z_2^{(1)}}{\partial w_{1,0}^{(1)}}$$

$$+ \frac{\partial L}{\partial o_2} \cdot \frac{\partial o_2}{\partial h_2^{(2)}} \cdot \frac{\partial h_2^{(2)}}{\partial z_2^{(2)}} \cdot \frac{\partial z_2^{(2)}}{\partial h_2^{(1)}} \cdot \frac{\partial h_2^{(1)}}{\partial z_2^{(1)}} \cdot \frac{\partial z_2^{(1)}}{\partial w_{1,0}^{(1)}}$$

$$= (o_1 - t_1) \cdot 1 \cdot 2 \cdot 1 \cdot 1 + (o_2 - t_2) \cdot 1 \cdot 0.5 \cdot -0.5 \cdot 1 \cdot 1$$

$$= -5 + 1 = -4.$$

$$w_{1,0}^{(1)} = w_{1,0}^{(1)} - 0.1 \times -4 = 0.8$$

10. The  $L_1$  regularization of the model parameter  $\mathbf{w} = [w_1, w_2, \dots, w_n]^T$  is defined as:

$$\|\mathbf{w}\|_1 = \sum_i^n |w_i|.$$

Below, we apply  $L_1$  regularization to a neural network model, of which the original objective function is defined as  $J(\mathbf{w}; X, y)$ , where  $\mathbf{w}$  is the model weights,  $X$  is the input, and  $y$  is the ground truth labels. The regularized objective function after adding the  $L_1$  normalization term becomes:

$$J_R(\mathbf{w}; X, y) = J(\mathbf{w}; X, y) + \alpha \|\mathbf{w}\|_1.$$

- (a) (2 points) Write the gradient  $\partial J_R / \partial \mathbf{w}$  in terms of  $\partial J / \partial \mathbf{w}$ .
- (b) (15 points) Finding the closed form solution for the root  $\mathbf{w}_R = [w_{R1}, w_{R2}, \dots, w_{Rn}]^T$  of  $\partial J_R / \partial \mathbf{w} = 0$  is difficult. Hence, we apply Taylor expansion to approximate  $J_R(\mathbf{w}; X, y)$ , and discard the high-order terms:

$$\hat{J}(\mathbf{w}; X, y) \approx J(\mathbf{w}^*) + \frac{1}{2}(\mathbf{w} - \mathbf{w}^*)^T H(\mathbf{w} - \mathbf{w}^*),$$

where  $H$  is the Hessian matrix, and  $\mathbf{w}^*$  is the solution of  $\partial J / \partial \mathbf{w} = 0$  (i.e., the optimal parameter for the objective function  $J$  without the regularization term). Suppose that it is known that  $H$  is a diagonal matrix with  $H = \text{diag}([H_{1,1}, \dots, H_{n,n}])$ , where  $H_{i,i} > 0, \forall i \in n$ . Write down a closed form expression for the root of  $\mathbf{w}_R$ .

(a).  $\frac{\partial J_R}{\partial \mathbf{w}} = \frac{\partial J}{\partial \mathbf{w}} + \lambda \mathbf{z} \cdot \text{sign}(\mathbf{w}).$

(b).

